



2016 IT Symposium: An Introduction to Docker



Bob Killen: Advanced Research Computing

Chris Kretler: ITS - Teaching and Learning





Please Do!

- Download or git clone: <u>https://github.com/chriskretler/nginx-docker-demo</u>
- 2. If using the zip download, unpack in a clean local directory
- 3. docker pull nginx





Who are we?

- Chris Kretler
- <u>ckretler@umich.edu</u>
- github.com/chriskretler
- With U since 2007
- Helping development groups on campus get to know and use Docker
- Currently leading ITS effort to deploy docker apps on OpenShift

- Bob Killen
- rkillen@umich.edu
- github.com/mrbobbytables
- @mrbobbytables
- University employee for 14 years
- Over 3 years of experience with Docker
- Deployed Docker in a production setting
- Experience with all major container orchestration platforms (Cattle, Kubernetes, Mesos and Swarm)





Overview

- Docker Overview
- Benefits of Using Docker
- Brief Tutorial
- More Features





What is Docker?



- An Open Source platform for developing, distributing, and running applications.
- Built around the concept of a 'container'.
 - Lightweight OS level virtual machine.
 - Bundles an application, its dependencies, and other run-time requirements into an immutable, redistributable image.
- Powered by Linux Container Subsystem
 - Provides isolation of resources CPU, memory, and network via kernel namespaces and cgroups.
- Available on multiple platforms and architectures
 - Linux (x86_64 and ARM)
 - Windows 10, server 2016
 - OSX 10.10+

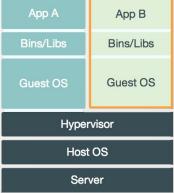




What Docker is not



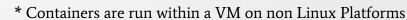
- A VM.
 - Docker is not a hypervisor, and a container is not a full virtual machine*
- Persistent.
 - A container's default behaviour is to be ephemeral.
- Secure by default.
 - A container is not a security panacea. It does have the capability of being more secure, but effort is required.



Virtual

Machine

App B Bins/Libs **Docker Engine Host OS** Server



Container





Docker Core Components



Engine	Image	Registry	
 The Docker Daemon running on a host. Manages building, storing, and running images on a specific host. 	 The item that is executed by the engine. Images consists of a manifest and collection of read-only layers generated by a Dockerfile. 	 The service that acts as an image repository. Example: Docker hub 	



Docker Workflow



Build	Ship	Run
 Build, develop and join the components of your application together into an image or set of images. 	 Push your image to a registry making it widely available. 	 Deploy, manage and use the same containers your application was developed on in a production environment.





Build - Images



- Docker images are built from a 'Dockerfile'.
- Dockerfiles contains a list of commands and actions that are used by the Docker Engine to assemble an image. Examples include 'FROM,' 'RUN', 'COPY', 'EXPOSE' and 'CMD'.

helloworld (demo1)

```
FROM alpine

RUN echo "hello world" > example

CMD ["cat", "example"]
```



Dockerfile Command Reference: https://docs.docker.com/engine/reference/builder/



Build - Images



helloworld (demol)

```
FROM alpine
RUN echo "hello world" > example
CMD ["cat", "example"]
```

- Use Alpine as the base image.
- Run the command writing 'hello world' to 'example'
- Instruct the container to execute the 'cat example' command when run.

```
muninn:demo1 bob$ docker build -t helloworld .

Sending build context to Docker daemon 2.048 kB

Step 1 : FROM alpine
---> baa5d63471ea

Step 2 : RUN echo "hello world" > example
---> Running in 931cf474a585
---> 3f2c16891e5b

Removing intermediate container 931cf474a585

Step 3 : CMD cat example
---> Running in b92112c29e03
---> 70289614bb95

Removing intermediate container b92112c29e03

Successfully built 70289614bb95
```



muninn:demo1 bob\$ docker run helloworld hello world



Build - Images - Layers



- An image is composed of layers.
- Each layer has an associated UUID and checksum.
- If multiple images do the same thing within their Dockerfiles, only a single layer for that step will be created.

helloworld (demo1)

layers (demo2)

```
FROM alpine
RUN echo "hello world" > example
CMD ["cat", "example"]
```

```
FROM alpine
RUN echo "hello world" > example
CMD [ "echo", "I'm in a container~!"]
```





Build - Images - Layers



```
muninn:demo1 bob$ docker history helloworld
IMAGE
                    CREATED
                                         CREATED BY
                                                                                         SIZE
                                         /bin/sh -c #(nop) CMD ["cat" "example"]
70289614bb95
                    About a minute ago
                                                                                         0 B
                    About a minute ago
                                        /bin/sh -c echo "hello world" > example
3f2c16891e5b
                                                                                         12 B
                                         /bin/sh -c #(nop) ADD file:7afbc23fda8b0b3872
baa5d63471ea
                    3 weeks ago
                                                                                         4.803 MB
```

baa5d3471ea 3f2c16891e5b

Base layers are the same.

muninn:demo2 bob	\$ docker history laye	rs	
IMAGE	CREATED	CREATED BY	SIZE
d3cb4c2894ff	4 seconds ago	/bin/sh -c #(nop) CMD ["echo" "I'm in a cont	0 B
3f2c16891e5b	34 minutes ago	/bin/sh -c echo "hello world" > example	12 B
baa5d63471ea	3 weeks ago	/bin/sh -c #(nop) ADD file:7afbc23fda8b0b3872	4.803 MB





Build - Images - FROM



- Images inherit the parent layers via the FROM directive.
- Encourages the workflow pattern of chaining images together.
- Docker Hub and many other build systems can trigger automatic rebuild of a container upon update of the parent.

from (demo3)

```
FROM layers
CMD [ "cat", "example"]
```

muninn:demo3 bob\$ docker run from hello world



```
CREATED
                                        CREATED BY
                                                                                         SIZE
                   46 minutes ago
                                        /bin/sh -c #(nop) CMD ["cat" "example"]
121b43b4f1e9
                                                                                        0 B
                    19 hours ago
                                        /bin/sh -c #(nop) CMD ["echo" "I'm in a cont
                                                                                        0 B
d3cb4c2894ff
                                        /bin/sh -c echo "hello world" > example
                    19 hours ago
                                                                                        12 B
f2c16891e5b
                                        /bin/sh -c #(nop) ADD file:7afbc23fda8b0b3872
 aa5d63471ea
                    3 weeks ago
```



Ship - Push



- A built image is tagged with a namespace and tag.
- A Namespace is usually a user or organization name combined with the name of the container itself.
- A tag is a little piece of metadata about the image, usually related to the version of the application in the image or the version of the image itself.
- Hub Name Format: <namespace>:<tag> mrbobbytables/layers:1
- Private Registry Name Format: <private registry url>/<namespace>:<tag>
 - registry.example.com/mrbobbytables/layers:1
- Only pushes layers that are not stored within the registry already.

```
muninn:demo2-layers bob$ docker tag layers mrbobbytables/layers:1
muninn:demo2-layers bob$ docker push mrbobbytables/layers:1
The push refers to a repository [docker.io/mrbobbytables/layers]
af7da5be3005: Layer already exists
011b303988d2: Layer already exists
1: digest: sha256:d8a2fefd97c0c2b86d2f2e6457ed5fb6c6d64c5017056e6145c47a7263b7a3a3 size: 735
muninn:demo2-layers bob$ docker pull mrbobbytables/layers:1
1: Pulling from mrbobbytables/layers

Digest: sha256:d8a2fefd97c0c2b86d2f2e6457ed5fb6c6d64c5017056e6145c47a7263b7a3a3
Status: Image is up to date for mrbobbytables/layers:1
```





Ship - Pull



Pulling an image only pulls layers not currently downloaded.

muninn2:demo2-layers bob\$ docker pull mrbobbytables/layers:1

1: Pulling from mrbobbytables/layers

Digest: sha256:d8a2fefd97c0c2b86d2f2e6457ed5fb6c6d64c5017056e6145c47a7263b7a3a3

Status: Downloaded newer image for mrbobbytables/layers:1





Run



- Upon run, containers are assigned a unique id and randomly generated name (unless supplied)
- Resources must be explicitly mapped or exposed (ports, volumes etc.)
- These resources can be shared with other containers, or exposed on the host itself.
- It's best to think of these exposed resources as a service, or the consumable portion of the container.





Run



muninn2:demo4-run bob\$ docker run -d -v \$(pwd):/usr/share/nginx/html -p 8888:80 nginx:stable-alpine 320a6676b00c995e3e1a47c2085219244e08a496bff0272a976ea53282e37415



• -d

tells docker to 'daemonize' the instance.

-v \$(pwd):/usr/share/nginx/html

 mounts the current working directory to the /usr/share/nginx/html directory within the container.

-p 8888:80

instructs the engine to map the host port
 8888 to the container port 80.

nginx:stable-alpine

The name of the image to run



Run



Docker Run Reference:

https://docs.docker.com/engine/reference/run/





Overview

- Docker Overview
- Benefits of Using Docker
- Brief Tutorial
- More Features





Benefits: Development

- Application environmental dependencies are bundled & specified:
 - No conflicts with existing applications in multi-tenanted situation.
 - What you build on your laptop will run the same way in any environment.
 - CI: Promotion with confidence.

```
UNIVERSITY OF MICHIGAN
```

```
FROM ubuntu: 16.04
MAINTAINER Teaching and Learning Devs <its-tl-dev@umich.edu>
# execute this separate to allow caching
RUN apt-get update
# install dependencies and link node to node;s, as ubuntu doesn't do this by default.
RUN apt-get install -v nodeis curl libmysglclient-dev python-dev gunicorn python-pip
django-filter python-ldap npm git libssl-dev libffi-dev xmlsec1 && \
    ln -s /usr/bin/nodeis /usr/bin/node
# install dependent python packages that aren't available as pre-built libraries
RUN pip install --upgrade pip && \
    pip install Django==1.9 cryptography==1.5.1 mysql-python djangorestframework
django-crispy-forms whitenoise requests djangosaml2 coverage
# install node packages that aren't available as pre-built libraries
RUN npm install -q bower grunt grunt-cli
# expose port, do this before source code install to minimize layering
EXPOSE 8000
# create place for app to run from
WORKDIR /app/
COPY . /app/
RUN echo '{ "allow_root": true }' > /root/.bowerrc
# Install Ocellus
RUN cd mbofui && \
    bower install && \
    nom install && \
    arunt docker
# copy settings file and launch django
CMD python manage.py migrate: gunicorn --workers=4 --bind=0.0.0.0:8000
hacks_mbof.wsqi:application
```



Benefits: Development

- Image registries (Docker Hub) facilitates experimentation with different platforms
 - Officially maintained environments for Python, Ruby, NodeJS, Java etc.
 - Components like redis, mongo, elasticsearch, postgress
 - latest / specific tags



redis officia	3.0K STAR:		DETAILS
busyl Busyl official	860 STAR:	10M+ S PULLS	> DETAILS
ubun officia	5.1K STAR:		> DETAILS
docker official	1.2K STAR:		DETAILS
alpino officia	1.6K STAR		DETAILS
mong officia	2.6K STAR		DETAILS
MySQL. officia	3.4K STAR:	100000	DETAILS
swam officia	543 STAR:		> DETAILS



Benefits: Operations

- Robust production ecosystem:
 - Zero-downtime deployments
 - Automatic scaling
 - Lower resource usage
 - Application separation
- Operations can maintain base images with organizational requirements bundled in private registries.



PINE ☆	
pushed: a month ago	
Scanned Images ②	
edge Compressed size; 2 MB Scanned a month ago	✓ This image has no known vulnerabilities
3.4 Compressed size: 2 MB Scanned a month ago	✓ This image has no known vulnerabilities
atest Compressed size: 2 MB Scanned 20 days ago	This image has no known vulnerabilities
3.3 Compressed size: 2 MB Scanned a month ago	✓ This image has no known vulnerabilities



Overview

- Docker Overview
- Benefits of Using Docker
- Brief Tutorial
- More Features





Demo

- You are going to have a working web application at the end of this demo!
- Download repo: https://github.com/chriskretler/nginx-docker-demo
- Application overview
- Docker build based on downloaded files
- Docker run & again
- Modify text blob index.htm
- Docker build to create image
 - Talk about caching and layers
- Docker run
 - See your updates





Overview

- Docker Overview
- Benefits of Using Docker
- Brief Tutorial
- More Features





Docker-Compose



- Compose is a lightweight tool built to define and run multi-container applications.
- It uses a simple, but feature complete yaml DSL to describe container infrastructure in a 'docker-compose.yml' file.
- It manages container infrastructure by breaking it down to 3 major components:
 - Services
 - Networks
 - Volumes







docker-compose.yml



compose-wordpress/docker-compose.yml

```
version: '2'
ervices:
 wordpress:
    depends_on:
      - db
   image: wordpress:latest
    networks:

    backend

    environment:
      WORDPRESS_DB_HOST: db:3306
      WORDPRESS DB PASSWORD: wordpress
db:
   image: mysal:5.7
   networks:

    backend

   volumes:
    - db_data:/var/lib/mysql
   environment:
    MYSOL_ROOT_PASSWORD: wordpress
    MYSQL_DATABASE: wordpress
    MYSQL_USER: wordpress
    MYSOL_PASSWORD: wordpress
 1b:
   image: dockercloud/haproxy
   depends_on:

    wordpress

   ports:
     - 80:80
   networks:

    backend

     - public
   environment:
     - ADDITIONAL_SERVICES=composewordpress:wordpress
     - /var/run/docker.sock:/var/run/docker.sock
```

```
networks:
   public:
    driver: bridge
   backend:
   driver: bridge
   ipom:
    driver: default
   config:
    - subnet: 192.168.223.0/24
    gateway: 192.168.223.1

volumes:
   db_data:
    driver: local
```

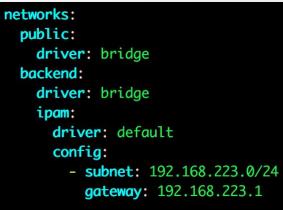




Compose-Networks



- Networks directive define consumable networks for the services.
- There are a variety of drivers, but the most common include:
 - Bridge (default)
 - Host
 - Overlay (multi-host capable)
- Containers can be attached to more than one network.
 - A 'frontend' or 'public' network will usually map a load balancer to one or more web servers.
 - A backend network will map the web servers to the database or application servers.







Compose-Volumes



The volumes directive contains named volumes only and does **NOT** include standard bind mounted volumes.

volumes:
 db_data:
 driver: local

- Volumes created this way do not map to a user-definable directory on the host unless using a 3rd party driver ('local-persist').
 - https://github.com/CWSpear/local-persist
- Intended purpose is to consume storage in a cluster deployment.
- For cluster deployments, evaluate EMC's rexray driver:
 - https://github.com/codedellemc/rexray





Compose-Services



- A service represents a container and defines the ways it interacts with external systems as well as other containers.
- Example: 'depends_on' configures startup order. The 'wordpress' container will not attempt to start until the 'db' container is up and running.
- Services can be attached to explicitly defined networks.
- Volumes in the service context can reference both bind mounts and named volumes.



```
wordpress:
  depends_on:
  image: wordpress:latest
  networks:
    - backend
  environment:
    WORDPRESS_DB_HOST: db:3306
    WORDPRESS_DB_PASSWORD: wordpress
 image: mysal:5.7
 networks:

    backend

 volumes:
   db_data:/var/lib/mysal
 environment:
   MYSOL ROOT PASSWORD: wordpress
   MYSQL_DATABASE: wordpress
   MYSOL_USER: wordpress
   MYSQL_PASSWORD: wordpress
 image: dockercloud/haproxy
 depends_on:

    wordpress

 ports:
   - 80:80
 networks:

    backend

   - public
 environment:

    ADDITIONAL_SERVICES=composewordpress:wordpress

    /var/run/docker.sock:/var/run/docker.sock
```



Orchestration















Orchestration



Kubernetes	Mesos	Rancher	Swarm
 Developed by Google Large Community Built for Web Scale Very flexible Complex High learning curve 	 Apache Project. Designed to be the 'kernel' of your datacenter Best for Data Science applications Complex High Learning Curve 	 Easy to use Can manage other orchestration systems Large catalog of services already built Large system overhead Not as fault tolerant as other systems 	 Easiest to deploy and use Extensive security built into engine Best when bringing a compose based stack Not as flexible as other options





<u>Where to go from here</u>



- Demo Files:
 - https://github.com/arc-ts/umitsm16-docker
- Docker Labs (Tutorials):
 - https://github.com/docker/labs
- Docker Run Reference:
 - https://docs.docker.com/engine/reference/run/
- Dockerfile Command Reference:
 - https://docs.docker.com/engine/reference/builder/
- Dockerfile Best Practices:
 - https://docs.docker.com/engine/userguide/eng-image/dockerfile_best-practices/
- Dockerfile Linter:
 - http://hadolint.lukasmartinelli.ch/





Where to go from here



Additional Questions

rkillen@umich.edu

ckretler@umich.edu

ARC-TS hosting Docker / Containerization workshops in the future.

http://arc-ts.umich.edu

